SHASHANK S

+91 94830 71522 | shashu.s08@gmail.com | LinkedIn | GitHub

PROFESSIONAL SUMMARY

GenAI Engineer with nearly 2 years of experience developing LLM-powered applications using Python and modern AI frameworks. Experienced in RAG systems, hybrid information retrieval, AI agents, and LLM evaluation, with hands-on experience integrating LLMs through FastAPI. Co-developed AgentFabric, an agentic AI platform and developed DevDocs, a RAG-based assistant platform for codebases.

TECHNICAL SKILLS

Languages: Python
GenAI and LLMs: OpenAI GPT models, Claude models, RAG, Hybrid retrieval, Chunking strategies, Embeddings, Tool and Function calling, AI agents and agentic workflows, Multi-agent orchestration, LLM evaluation
Frameworks and Libraries: LangChain, LangGraph, FastAPI
Databases: MySQL, PostgreSQL, Vector databases (Qdrant, pgvector, Elasticsearch)
Cloud and Tools: Azure, Docker, Git, GitHub
AI-Assisted Development: GitHub Copilot, Claude Code, Codex

PROFESSIONAL EXPERIENCE

Wipro Limited , Bengaluru March, 2025 – Present
GenAI Engineer

RAG Platform for a Banking Client

- Developed an end-to-end RAG platform covering policy document ingestion, chunking, indexing and embedding.
- Implemented hybrid retrieval combining BM25 and vector search (cosine similarity), merged with Reciprocal Rank Fusion and refined with a reranker.
- Delivered streamed LLM answers (using OpenAI GPT model) with citations to the source files used.

WEGA-Studio Application

- Co-developed an internal application that uses Factory AI's Droid Exec to run agentic workflows across multiple use cases.
- Workflows include documentation generation, user story generation with push to Jira via MCP, implementation of user stories as code features using Jira and Git MCP, and test case generation for Git repositories via Git MCP.

TECHNICAL PROJECTS

DevDocs — RAG Assistant for Codebases | *Python*, *FastAPI*, *React*, *PostgreSQL*, *Elasticsearch*, *Docker* September, 2026

- Built a hybrid retrieval pipeline combining Elasticsearch BM25 and vector search using Reciprocal Rank Fusion (RRF), with similarity-based gating to reject off-topic queries before LLM invocation.
- Implemented incremental ingestion using SHA-256 change detection, structure-aware chunking (Python AST, Markdown, PDF), and selective re-embedding, with PostgreSQL as the source of truth.
- Enforced project isolation and generated file-and-line citations in application code to ensure grounded, verifiable answers.
- Implemented streaming chat, conversation memory, Stop/Retry controls, and fallback-model routing, reducing burst-question latency from 40+ seconds to approximately 2 seconds.
- Evaluated on a 428-question benchmark across a 104-file project, improving Recall@5 from 0.83 to 0.87 and MRR from 0.78 to 0.83 on tuning data, with further gains on held-out evaluation.

EDUCATION

SJB Institute of Technology 2020 – 2024
B.E, Computer Science and Technology CGPA: 8.93